

Lecture 4, 5, 6, 7: From Linear Perspectives and μP to the Muon Optimizer

Jincheng Ou

Fall 25

Contents

1 The Linear Perspective and Lazy Training	2
1.1 Taylor Expansion of Neural Networks	2
2 The General Optimizer Recipe: Constrained Optimization	2
2.1 The Objective Function	2
2.2 Derivation 1: The Infinity Norm (ℓ_∞) → SignSGD	2
2.3 Derivation 2: The Euclidean Norm (ℓ_2) → Standard GD	3
3 Initialization and Data Standardization	3
3.1 The "ReLU Corner" Problem	3
3.2 Xavier (Glorot) Initialization	4
4 Matrix Optimization: Shampoo and Semi-Orthogonality	4
4.1 The Matrix Trace Optimization	4
5 Maximal Update Parameterization (μP)	5
5.1 The RMS-RMS Norm	5
5.2 Scaling Laws for Feature Learning	5
6 Muon: Momentum Orthogonalized by Newton-Schulz	5
6.1 Key Idea 1: RMS-RMS Optimization Geometry	6
6.2 Key Idea 2: Newton-Schulz Iteration	6
6.3 The Muon Algorithm	6
7 Summary: The Evolution of Optimizers	7

1 The Linear Perspective and Lazy Training

1.1 Taylor Expansion of Neural Networks

We begin by treating a neural network $f(x, \theta)$ not as a black box, but as a function of its parameters $\theta \in \mathbb{R}^m$ for a fixed input x . We can analyze the training dynamics by linearizing the network around the current parameters θ_0 using a first-order Taylor expansion:

$$f(x, \theta_0 + \Delta\theta) \approx f(x, \theta_0) + \langle \nabla_\theta f(x, \theta_0), \Delta\theta \rangle + O(\|\Delta\theta\|^2) \quad (1)$$

Here, $\nabla_\theta f(x, \theta)$ acts as the "feature vector" of the model.

Definition 1.1: Lazy Training Regime

The **Lazy Training** assumption posits that during training, the weights θ move very little from their initialization θ_0 (i.e., $\|\Delta\theta\|$ is small). Consequently, the model behaves essentially as a linear model over fixed random features defined by $\nabla_\theta f(x, \theta_0)$.

This creates a fundamental tension in optimization:

- **Desire for Speed:** We want large updates ($\Delta\theta$) to minimize loss quickly.
- **Validity of Approximation:** We need small updates to ensure the Taylor approximation holds (trust region).

2 The General Optimizer Recipe: Constrained Optimization

To resolve the tension above, we formalize optimization as minimizing the linearized loss subject to a geometry-aware constraint on the step size.

2.1 The Objective Function

We aim to find the update step $\Delta\theta$ that minimizes the change in loss \mathcal{L} :

$$\Delta\theta^* = \underset{\Delta\theta}{\operatorname{argmin}} \langle \nabla_\theta \mathcal{L}(\theta), \Delta\theta \rangle \quad \text{s.t.} \quad \|\Delta\theta\|_P \leq \eta \quad (2)$$

where $\|\cdot\|_P$ is a chosen norm defining the geometry of the trust region, and η is the step size.

2.2 Derivation 1: The Infinity Norm (ℓ_∞) → SignSGD

Let the trust region be defined by the ℓ_∞ norm: $\|\Delta\theta\|_\infty \leq \eta$. This implies $\max_j |\Delta\theta[j]| \leq \eta$. The problem decouples into coordinate-wise optimization:

$$\min_{\Delta\theta} \sum_j \nabla_\theta \mathcal{L}[j] \cdot \Delta\theta[j] \quad \text{s.t.} \quad -\eta \leq \Delta\theta[j] \leq \eta \quad \forall j \quad (3)$$

Derivation 2.1: Optimal Solution for ℓ_∞

Since the objective is linear, the minimum occurs at the boundaries.

- If $\nabla\mathcal{L}[j] > 0$, we need $\Delta\theta[j]$ to be negative to minimize the product. We choose the largest magnitude allowed: $-\eta$.
- If $\nabla\mathcal{L}[j] < 0$, we choose $+\eta$.

Thus:

$$\Delta\theta[j] = -\eta \cdot \operatorname{sign}(\nabla\mathcal{L}[j]) \quad (4)$$

This recovers SignSGD, which is the geometric foundation of Adam. Adam essentially performs SignSGD with momentum and de-biasing.

2.3 Derivation 2: The Euclidean Norm (ℓ_2) → Standard GD

Let the constraint be $\|\Delta\theta\|_2 \leq \eta$. We minimize $\langle g, \Delta\theta \rangle$. By the Cauchy-Schwarz inequality:

$$|\langle g, \Delta\theta \rangle| \leq \|g\|_2 \|\Delta\theta\|_2$$

The inner product is minimal (most negative) when $\Delta\theta$ is exactly **anti-aligned** with the gradient g :

$$\Delta\theta = -\eta \frac{g}{\|g\|_2} \quad (5)$$

This corresponds to normalized Gradient Descent.

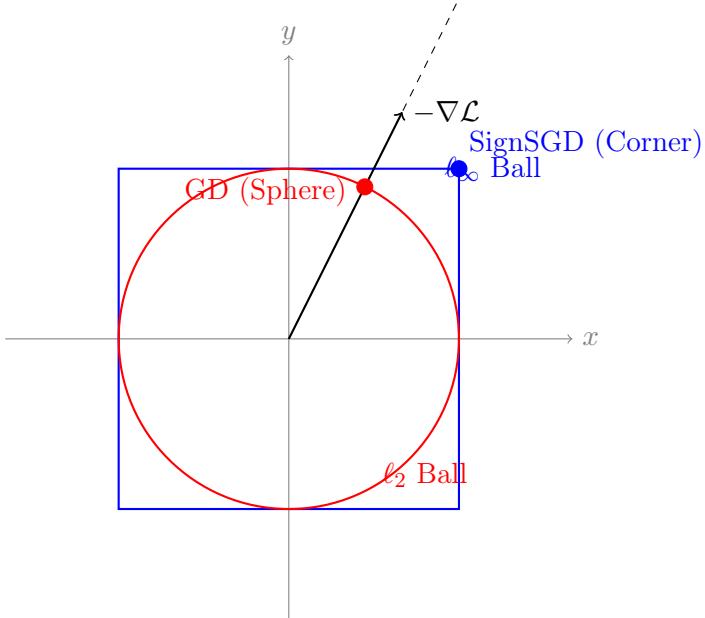


Figure 1: Geometric interpretation of Optimizers. SignSGD moves to the corners of the hypercube (using component-wise max steps), while GD moves to the boundary of the hypersphere.

3 Initialization and Data Standardization

Effective optimization requires that the inputs to non-linearities fall within their "active" regions.

3.1 The "ReLU Corner" Problem

Consider a ReLU unit $y = \max(0, wx + b)$. The "corner" or "elbow" occurs at $x = -b/w$.

- If input data is unstandardized (e.g., range [100, 200]) and initialization is standard normal $\mathcal{N}(0, 1)$, the probability of the "elbow" falling within the data range is nearly zero ($\approx 0.1\%$).
- **Consequence:** The neuron behaves purely linearly or outputs zero for all data. Feature learning (which relies on non-linearity) fails.

This necessitates **data standardization** (mean 0, variance 1).

3.2 Xavier (Glorot) Initialization

We aim to preserve the variance of activations across layers to prevent vanishing/exploding gradients. Consider a linear layer $h_l = Wh_{l-1}$ where $W \in \mathbb{R}^{d_{out} \times d_{in}}$. Assuming inputs are independent with 0 mean:

$$\text{Var}(h_l) = \text{Var}\left(\sum_{j=1}^{d_{in}} W_{ij} h_{l-1,j}\right) \quad (6)$$

$$= \sum_{j=1}^{d_{in}} \text{Var}(W_{ij} h_{l-1,j}) \quad (\text{Independence}) \quad (7)$$

$$= \sum_{j=1}^{d_{in}} \mathbb{E}[W_{ij}^2] \mathbb{E}[h_{l-1,j}^2] \quad (\text{Zero Mean}) \quad (8)$$

$$= d_{in} \cdot \text{Var}(W) \cdot \text{Var}(h_{l-1}) \quad (9)$$

To preserve variance ($\text{Var}(h_l) = \text{Var}(h_{l-1})$), we require:

$$d_{in} \cdot \text{Var}(W) = 1 \implies \text{Var}(W) = \frac{1}{d_{in}} \quad (10)$$

This derivation motivates the **RMS Norm** used in later optimizers.

4 Matrix Optimization: Shampoo and Semi-Orthogonality

Neural networks operate on matrices, not flat vectors. Optimizing in the matrix space yields better convergence properties.

4.1 The Matrix Trace Optimization

We reformulate the optimizer recipe for a weight matrix W . Objective: Maximize the correlation with the gradient $G = \nabla_W \mathcal{L}$ subject to a Spectral Norm constraint.

$$\max_{\Delta W} \text{Trace}(G^\top(-\Delta W)) \quad \text{s.t.} \quad \|\Delta W\|_{\text{spectral}} \leq \eta \quad (11)$$

Derivation 4.1: Deriving the Shampoo Update

Let the Singular Value Decomposition (SVD) of the gradient be $G = U\Sigma V^\top$. Substitute into the trace:

$$\begin{aligned} \text{Trace}(G^\top(-\Delta W)) &= \text{Trace}((U\Sigma V^\top)^\top(-\Delta W)) \\ &= \text{Trace}(V\Sigma U^\top(-\Delta W)) \\ &= \text{Trace}(\Sigma U^\top(-\Delta W)V) \quad (\text{Cyclic property}) \end{aligned}$$

Let $Z = U^\top(-\Delta W)V$. The problem becomes maximizing $\text{Trace}(\Sigma Z)$ subject to $\|Z\|_2 \leq \eta$ (since unitary matrices U, V don't change spectral norm).

$$\text{Trace}(\Sigma Z) = \sum_i \sigma_i Z_{ii}$$

Since $\sigma_i \geq 0$, we maximize this by making Z diagonal with entries equal to the max allowed value η .

$$Z = \eta I \implies U^\top(-\Delta W)V = \eta I \implies \Delta W = -\eta UV^\top$$

Shampoo Implementation: Computing SVD is expensive (Shampoo approximates this using 4th roots of

$$W_{t+1} = W_t - \eta L_t^{-1/4} G_t R_t^{-1/4}$$

where $L_t \approx GG^\top$ and $R_t \approx G^\top G$.

5 Maximal Update Parameterization (μP)

Standard optimizers (like Adam) often require re-tuning the Learning Rate (LR) when the model width changes. μP provides a scaling law to transfer hyperparameters from small to large models.

5.1 The RMS-RMS Norm

Motivated by Xavier initialization, we define a matrix norm that is invariant to matrix .

Definition 5.1: Induced RMS-RMS Norm

$$\|A\|_{RMS \rightarrow RMS} = \max_{\|x\|_{RMS}=1} \|Ax\|_{RMS}$$

where $\|x\|_{RMS} = \frac{1}{\sqrt{d_{in}}} \|x\|_2$.

Theorem 5.1: Relation to Spectral Norm

The RMS-RMS norm is a scaled version of the spectral norm:

$$\|A\|_{RMS \rightarrow RMS} = \sqrt{\frac{d_{in}}{d_{out}}} \|A\|_2 \quad (12)$$

5.2 Scaling Laws for Feature Learning

For a network to learn features effectively as width increases, two conditions must hold:

1. **Activations are O(1):** $\|h_l\|_{RMS} = \Theta(1)$.
2. **Updates are O(1):** $\|\Delta h_l\|_{RMS} = \Theta(1)$.

Using $\Delta h = \Delta Wh$:

$$\|\Delta h\|_{RMS} \leq \|\Delta W\|_{RMS \rightarrow RMS} \|h\|_{RMS}$$

Since $\|h\| \approx 1$, we need $\|\Delta W\|_{RMS \rightarrow RMS} \approx 1$. Converting this back to the standard spectral norm updates used by optimizers:

$$\|\Delta W\|_2 \leq \eta \sqrt{\frac{d_{out}}{d_{in}}} \quad (13)$$

This implies the optimal learning rate should scale with $\sqrt{\frac{d_{out}}{d_{in}}}$ or inversely with fan-in d_{in} depending on the specific parameterization.

6 Muon: Momentum Orthogonalized by Newton-Schulz

Muon is an optimizer designed to combine the spectral efficiency of Shampoo with the low computational cost required for large-scale training (like LLMs).

6.1 Key Idea 1: RMS-RMS Optimization Geometry

Muon modifies the optimizer recipe to use the RMS-RMS norm constraint instead of the Spectral norm.

$$\operatorname{argmin} \langle G, \Delta W \rangle \quad \text{s.t.} \quad \|\Delta W\|_{RMS \rightarrow RMS} \leq \eta$$

Using the scaling relation derived in Theorem 9.1, the update becomes:

$$\Delta W^* = -\eta \sqrt{\frac{d_{out}}{d_{in}}} U V^\top \quad (14)$$

This automatically injects the architecture-dependent scaling factors (μP) into the update.

6.2 Key Idea 2: Newton-Schulz Iteration

Instead of computing SVD (UV^\top) which is slow, Muon approximates the semi-orthogonal matrix UV^\top using an iterative polynomial method.

Theorem 6.1: Newton-Schulz Convergence

The iteration $X_{k+1} = \frac{3}{2}X_k - \frac{1}{2}X_k X_k^\top X_k$ drives singular values towards 1.

Derivation 6.1: Polynomial SVD Interaction

Since odd polynomials commute with SVD, applying polynomial $P(X)$ to matrix $X = U\Sigma V^\top$ is equivalent to applying $P(\sigma)$ to the singular values [cite: 76-78]. Consider $P(x) = 1.5x - 0.5x^3$.

- If $\sigma = 1$: $P(1) = 1.5(1) - 0.5(1) = 1$ (Stable fixed point).
- If σ is small: $P(\sigma) \approx 1.5\sigma$ (Amplification).

Iterating this pushes singular values $\sigma \in (0, \sqrt{3})$ towards 1.

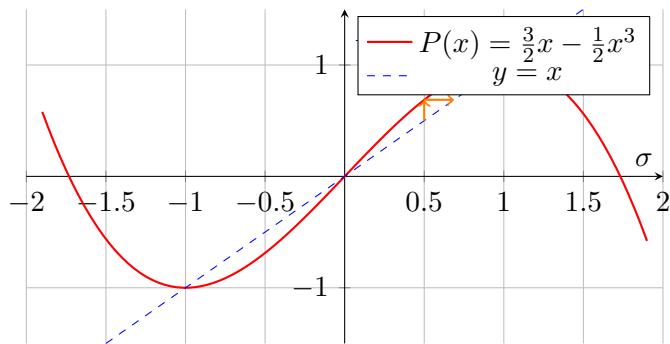


Figure 2: Newton-Schulz Polynomial. Values in $(0, \sqrt{3})$ converge to 1.

6.3 The Muon Algorithm

Muon replaces the adaptive arithmetic of Adam with Newton-Schulz orthogonalization:

Muon Update Step

1. **Momentum:** $B_t = \mu B_{t-1} + \nabla \mathcal{L}_t$
2. **Normalization:** To ensure convergence of Newton-Schulz (singular values $< \sqrt{3}$), normalize by Frobenius norm:

$$\hat{B}_t = \frac{B_t}{\max(1, \|B_t\|_F / \sqrt{d_{in} d_{out}})}$$

3. **Newton-Schulz (approx 5 iterations):**

$$X_0 = \hat{B}_t, \quad X_{k+1} = \frac{3}{2}X_k - \frac{1}{2}X_k X_k^\top X_k$$

4. **Parameter Update:**

$$W_{t+1} = W_t - \eta \cdot \sqrt{\frac{\max(1, d_{out})}{\max(1, d_{in})}} \cdot X_{\text{final}}$$

7 Summary: The Evolution of Optimizers

1. **SGD:** Linear ℓ_2 constraint \rightarrow Scale by gradient magnitude.
2. **Adam/SignSGD:** Linear ℓ_∞ constraint \rightarrow Scale by coordinate-wise magnitude (sign).
3. **Shampoo:** Matrix Spectral constraint \rightarrow Scale by singular values (Preconditioners).
4. **Muon:** Matrix RMS-RMS constraint + Newton-Schulz \rightarrow Architecture-aware scaling + Efficient semi-orthogonalization.